


The Case for Tiny Tasks in Compute Clusters

Kay Ousterhout*, Aurojit Panda*, Joshua Rosen*,
Shivaram Venkataraman*, Reynold Xin*,
Sylvia Ratnasamy*, Scott Shenker*+, Ion Stoica*



* UC Berkeley, + ICSI

Setting

Map Reduce/Spark/Dryad
Job




Use smaller tasks!




Why? How? Where?

Why? How? Where?

Problem: Skew and Stragglers




Benefit: Handling of Skew and Stragglers





As much as 5.2x reduction in job completion time!

Problem: Batch and Interactive Sharing

Clusters forced to trade off utilization and responsiveness!




Benefit: Improved Sharing



High-priority tasks not subject to long wait times!

Benefits: Recap


(1) Straggler mitigation



Mantri (OSDI '10)
Scarlett (EuroSys '11)
SkewTune (SIGMOD '12)
Dolly (NSDI '13)

...

(2) Improved sharing



Quincy (SOSP '09)
Amoeba (SOCC '12)

...


Why? How? Where?

Schedule task

Scheduling requirements:

High Throughput (millions per second)

Low Latency (milliseconds)



Distributed Scheduling
(e.g., Sparrow Scheduler)

Schedule
task

Launch
task

Use existing thread pool to
launch tasks


Schedule
task

Launch
task

Use existing thread pool to
launch tasks

+

Cache task binaries



Task launch = RPC time (<1ms)

Schedule

task



Launch

task




Read input

data


Smallest efficient file
block size:

8MB



Distribute Metadata
(à la Flat Datacenter Storage, OSDI '12)

Schedule task
↓
Launch task
↓
Read input data
↓
Execute task + read data for next task



Tons of tiny transfers!
↓
Framework-Controlled I/O
(enables optimizations, e.g., pipelining)

Schedule
task



Launch
task

Read input
data



Execute task
+ read data
for next task

How low can you go?


8MB disk block





100's of
milliseconds

Why? How? Where?


Original Job




Tiny Tasks Job



Reduce Tasks



Original Reduce Phase



Tiny Tasks = ?


Splitting Large Tasks

- Aggregation trees
 - Works for functions that are associative and commutative
- Framework-managed temporary state store
- Ultimately, need to allow a small number of large tasks

Tiny tasks
mitigate stragglers

+

Improve sharing



Distributed scheduling → Launch task in existing thread pool → Distributed file metadata → Pipelined task execution

Questions? Find me or Shivaram:

